

tu11_re_PandasReview

February 23, 2023

1 Pandas Review

Pandas is a Python package for organizing and analyzing data. In one sense, it is a generalization of NumPy, on which it is based.

NumPy is fantastic for working with numerical data that are “well behaved”. For example, if you are analyzing data from a tightly controlled laboratory experiment, then NumPy might be perfect.

In the broader world of behavioral data science, however, data can be complicated. Variables can be of multiple types, values can be missing, etc. Pandas was developed to make it easier for us to work with data sets in general, not just numerical arrays.

If you have experience in R then, in a nutshell, pandas gives you an equivalent to R in Python (some data scientists use both, picking one or the other depending on the project, but most people prefer sticking with one language if they can).

1.1 Pandas Data

The main data object in pandas is the **DataFrame**. It is a table of data in which each column has a name, generally corresponding to a specific real-world variable.

Just as we can think about a NumPy array as a spatial layout of a Python list of lists, we can think of a pandas **DataFrame** as a spatial layout of a Python dictionary.

Consider the following Python dictionary:

```
[1]: dis_chars = {'name': ['Mickey', 'Minnie', 'Pluto'],  
                 'gender': ['m', 'f', 'n'],  
                 'age': [95, 95, 93],}
```

```
[2]: dis_chars
```

```
[2]: {'name': ['Mickey', 'Minnie', 'Pluto'],  
      'gender': ['m', 'f', 'n'],  
      'age': [95, 95, 93]}
```

On the one hand, this is a nice organized *container* of data. But on the other hand, it is not much else. If we wanted to compute anything, like the mean age of all non-male characters, we’d have to start writing code from scratch.

Let’s make our dictionary into a **DataFrame**. First, we’ll import pandas.

```
[3]: import pandas as pd
```

Importing pandas as `pd` is conventional, like importing numpy as `np`, so there's no reason to do anything else.

Now we can convert our data to a `DataFrame` using `pd.DataFrame()`.

```
[4]: dis_df = pd.DataFrame(dis_chars)
```

And let's look at our new creation!

```
[5]: dis_df
```

```
[5]:      name gender  age
0  Mickey      m   95
1  Minnie      f   95
2   Pluto      n   93
```

Now we have a nice organized table of data, in which each column corresponds to a variable, and can be referred to by name.

```
[6]: dis_df['name']
```

```
[6]: 0    Mickey
     1    Minnie
     2     Pluto
     Name: name, dtype: object
```

Further, it makes it relatively easy for us to do lots of analyses “out of the box”. For example:

```
[7]: dis_df['age'].mean()
```

```
[7]: 94.33333333333333
```

Here, we just grabbed a column of data by name (`dis_df['age']`), and then computed its mean with the built-in `mean()` method.

The `DataFrame` isn't the only type of object in pandas, but it's the biggie. If you have experience in R, then you'll be in familiar territory, because the `DataFrame` in Python is modeled after the data frame (or tibble) in R.

```
[8]: type(dis_df)
```

```
[8]: pandas.core.frame.DataFrame
```

Each column of a `DataFrame` is a pandas `Series`.

```
[9]: dis_age_s = dis_df['age']
     dis_age_s
```

```
[9]: 0    95
      1    95
      2    93
      Name: age, dtype: int64
```

```
[10]: type(dis_age_s)
```

```
[10]: pandas.core.series.Series
```

And each series is a collection of more fundamental objects. So if we look at the last age in our series...

```
[11]: a = dis_age_s[2]
      a
```

```
[11]: 93
```

And check the type...

```
[12]: type(a)
```

```
[12]: numpy.int64
```

We see that it is a numpy integer; a hint that pandas is indeed built from NumPy!

If we check the type of one of the other values:

```
[50]: type(dis_df['gender'][2])
```

```
[50]: str
```

We see that it is a Python string object. (Take a moment to dissect that line of code, and see how it is doing exactly the same thing as we did to get the type of an age value, just in one go.)

In the code cell below, get the very first name in our Disney DataFrame.

```
[15]: # At first, Mickey's name was going to be Mortimer Mouse. I know, right?
      dis_df['name'][0]
```

```
[15]: 'Mickey'
```

One great thing about pandas is that, if we want to add a column, we just act like it already exists and assign values to it. Like this:

```
[16]: dis_df['wearsBow'] = [False, True, False]
      dis_df
```

```
[16]:      name gender  age  wearsBow
0  Mickey      m   95     False
1  Minnie      f   95      True
2   Pluto      n   93     False
```

Notice that we are addressing a ‘wearsBow’ column just like we would an existing column such as ‘name’. Pandas, rather than complain and be annoying, just creates the column for us!

1.2 Data i/o (Input and Output)

One of the really great things about pandas is that it makes reading, inspecting, and writing data files in common formats very easy.

1.2.1 Importing (input)

Following the pandas documentation, let’s look at some data about the passengers on the RMS Titanic.

Download the `titanic.csv` and place in folder named ‘data’ that is in the same folder as you have this notebook.

Now, loading it is as easy as calling `pd.read_csv()`:

```
[20]: In [2]: titanic = pd.read_csv("data/titanic.csv")
```

There are lots of other formats that pandas can read, including excel and html.

It can even read data from the clipboard! Try it! Go to the [Wikipedia page for Austin](#), scroll to the demographics section, and select the three columns (including the headers) down to 2020, and copy them to your clipboard.

Now run the code below.

```
[21]: atx_pop = pd.read_clipboard()
```

```
[22]: atx_pop
```

```
[22]:      Historical population
Census Pop.      Note      %±
1850    629          -      None
1860   3,494    455.5%      None
1870   4,428    26.7%      None
1880  11,013   148.7%      None
1890  14,575    32.3%      None
1900  22,258    52.7%      None
1910  29,860    34.2%      None
1920  34,876    16.8%      None
1930  53,120    52.3%      None
1940  87,930    65.5%      None
1950 132,459    50.6%      None
1960 186,545    40.8%      None
```

1970	253,539	35.9%	None
1980	345,890	36.4%	None
1990	465,622	34.6%	None
2000	656,562	41.0%	None
2010	790,390	20.4%	None
2020	961,855	21.7%	None

1.2.2 Inspecting

It's important to peek at any imported data to make sure nothing looks funny (like we just did with the Austin population data). So let's peek at the RMS Titanic data.

[23]: `titanic`

```
[23]:      PassengerId  Survived  Pclass  \
0             1         0         3
1             2         1         1
2             3         1         3
3             4         1         1
4             5         0         3
..          ...         ...         ...
886          887         0         2
887          888         1         1
888          889         0         3
889          890         1         1
890          891         0         3
```

```

                                Name      Sex  Age  SibSp  \
0                Braund, Mr. Owen Harris   male  22.0      1
1  Cumings, Mrs. John Bradley (Florence Briggs Th... female  38.0      1
2                Heikkinen, Miss. Laina   female  26.0      0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)   female  35.0      1
4                Allen, Mr. William Henry   male  35.0      0
..          ...         ...         ...         ...
886                Montvila, Rev. Juozas   male  27.0      0
887                Graham, Miss. Margaret Edith   female  19.0      0
888    Johnston, Miss. Catherine Helen "Carrie"   female   NaN      1
889                Behr, Mr. Karl Howell   male  26.0      0
890                Dooley, Mr. Patrick   male  32.0      0
```

```

      Parch      Ticket    Fare Cabin Embarked
0         0    A/5 21171   7.2500   NaN        S
1         0      PC 17599  71.2833   C85        C
2         0  STON/O2. 3101282   7.9250   NaN        S
3         0      113803  53.1000  C123        S
4         0      373450   8.0500   NaN        S
..      ...         ...         ...         ...
```

886	0	211536	13.0000	NaN	S
887	0	112053	30.0000	B42	S
888	2	W./C. 6607	23.4500	NaN	S
889	0	111369	30.0000	C148	C
890	0	370376	7.7500	NaN	Q

[891 rows x 12 columns]

A nice thing about pandas `DataFrames` is that, by default, they show you their first and last 5 rows (their head and tail), and then tell you how big they are (891x12 in this case).

We can look at as much of the head or tail as we want with the `head()` and `tail()` methods.

```
[24]: titanic.tail(9)
```

```
[24]:
```

	PassengerId	Survived	Pclass	Name \
882	883	0	3	Dahlberg, Miss. Gerda Ulrika
883	884	0	2	Banfield, Mr. Frederick James
884	885	0	3	Sutehall, Mr. Henry Jr
885	886	0	3	Rice, Mrs. William (Margaret Norton)
886	887	0	2	Montvila, Rev. Juozas
887	888	1	1	Graham, Miss. Margaret Edith
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"
889	890	1	1	Behr, Mr. Karl Howell
890	891	0	3	Dooley, Mr. Patrick

	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
882	female	22.0	0	0	7552	10.5167	NaN	S
883	male	28.0	0	0	C.A./SOTON 34068	10.5000	NaN	S
884	male	25.0	0	0	SOTON/OQ 392076	7.0500	NaN	S
885	female	39.0	0	5	382652	29.1250	NaN	Q
886	male	27.0	0	0	211536	13.0000	NaN	S
887	female	19.0	0	0	112053	30.0000	B42	S
888	female	NaN	1	2	W./C. 6607	23.4500	NaN	S
889	male	26.0	0	0	111369	30.0000	C148	C
890	male	32.0	0	0	370376	7.7500	NaN	Q

Use the cell below to display the first 11 rows of the titanic data.

```
[26]: # but these rows go to 11...
titanic.head(12)
```

```
[26]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	

4	5	0	3
5	6	0	3
6	7	0	1
7	8	0	3
8	9	1	3
9	10	1	2
10	11	1	3
11	12	1	1

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	
5	Moran, Mr. James	male	NaN	0	
6	McCarthy, Mr. Timothy J	male	54.0	0	
7	Palsson, Master. Gosta Leonard	male	2.0	3	
8	Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)	female	27.0	0	
9	Nasser, Mrs. Nicholas (Adele Achem)	female	14.0	1	
10	Sandstrom, Miss. Marguerite Rut	female	4.0	1	
11	Bonnell, Miss. Elizabeth	female	58.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S
5	0	330877	8.4583	NaN	Q
6	0	17463	51.8625	E46	S
7	1	349909	21.0750	NaN	S
8	2	347742	11.1333	NaN	S
9	0	237736	30.0708	NaN	C
10	1	PP 9549	16.7000	G6	S
11	0	113783	26.5500	C103	S

We can also look at the data types:

```
[27]: titanic.dtypes
```

```
[27]: PassengerId    int64
      Survived       int64
      Pclass         int64
      Name           object
      Sex            object
```

```

Age          float64
SibSp        int64
Parch        int64
Ticket       object
Fare         float64
Cabin        object
Embarked     object
dtype: object

```

(the columns listed as “object” seem to be strings)

We can also get more detailed information using the `info()` method:

```
[28]: titanic.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null   int64
1   Survived        891 non-null   int64
2   Pclass          891 non-null   int64
3   Name            891 non-null   object
4   Sex             891 non-null   object
5   Age             714 non-null   float64
6   SibSp           891 non-null   int64
7   Parch           891 non-null   int64
8   Ticket          891 non-null   object
9   Fare            891 non-null   float64
10  Cabin           204 non-null   object
11  Embarked        889 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB

```

This gives us a nice summary of the types of data in the columns and, in particular, how many valid (non-missing) values are in each. We can see that “Cabin”, for example, has many missing values.

1.2.3 Exporting (output)

The `to_` methods, such as `to_csv()`, `to_excel()`, etc., allow us to export data in many ways. As an example, let’s export the titanic data as a Microsoft Excel file.

In the cell below, use `titanic.to_excel(...)` to export the data to an Excel spreadsheet.

```
[51]: # exporting an Excel file!
titanic.to_excel('titanic.xlsx')
```


Open the file in Excel to verify that the export worked.

1.3 Selecting Data

In numpy, we select data by primarily by row and column indexes. In pandas, we generally address columns (corresponding to real world variables) by *name* and rows by one or more *criteria*.

1.3.1 Getting columns

As we did above with our little toy Disney data, we can compute the mean age of the passengers by grabbing that column of data by name, and then computing the mean of it.

```
[31]: ages = titanic['Age']  
      ages.mean()
```

```
[31]: 29.69911764705882
```

In the cell below, compute the mean age in one line of code (i.e., not creating the temporary ‘age’ object).

```
[47]: # average age of passengers on the RMS Titanic  
      titanic['Age'].mean()
```

```
[47]: 29.69911764705882
```

We can get multiple columns by indexing our `DataFrame` with a Python list of column names. We can do this in two lines for readability.

```
[39]: wanted_cols = ['Fare', 'Survived']  
      fare_surv = titanic[wanted_cols]
```

```
[40]: fare_surv
```

```
[40]:
```

	Fare	Survived
0	7.2500	0
1	71.2833	1
2	7.9250	1
3	53.1000	1
4	8.0500	0
..
886	13.0000	0
887	30.0000	1
888	23.4500	0
889	30.0000	1

```
890    7.7500    0
```

```
[891 rows x 2 columns]
```

But more commonly we do it in a single line.

```
[41]: fare_surv = titanic[['Fare', 'Survived']]
```

Your initial reaction might be “Why the double brackets? Why not single brackets?”, and the reason should be clear if we look back at the two line example: the `DataFrame` expects a Python list, not separate strings. So the outer set of brackets are indexing brackets, and the inner set defines a Python list.

1.3.2 Getting rows

We generally extract rows of interest by placing one or more criteria on a particular column.

```
[42]: my_critereon = fare_surv['Fare'] > 20
      rich = fare_surv[my_critereon]
```

What is actually happening here is that the logical test `fare_surv['Fare'] > 20` is creating a pandas series that is `True` for the rows in which the fare paid was greater than 20 pounds sterling, and `False` otherwise.

Let's look at `my_critereon`:

```
[43]: my_critereon
```

```
[43]: 0      False
      1       True
      2      False
      3       True
      4      False
      ...
      886     False
      887      True
      888      True
      889      True
      890     False
      Name: Fare, Length: 891, dtype: bool
```

This series is then used to get all the rows of `fare_surv` that correspond to the `True` values, and these are placed in `rich`.

This is known as *logical indexing*, and is widely used in data analysis!

As with fetching columns, we can do this one line instead of two.

```
[44]: rich = fare_surv[fare_surv['Fare'] > 20]
```

Whether you make a separate indexing series like `my_critereon` or put the test inside the indexing brackets is up to you. For simple tests, putting the test inside the brackets doesn't hurt the readability of the code at all. For more complicated tests – if you wanted all the cases of female passengers that paid between 20 and 50 lbs. for their fare, and had no siblings and two parents aboard, say – then you might want to make the test series first, and then do the indexing.

In the cell below, get the passenger class (`Pclass`) and survival status of passengers that paid more than 20 pounds for their voyage.

```
[57]: # passenger class and survival of high fares
pclass_surv = titanic[['Fare', 'Survived', 'Pclass']]
rich_pass = pclass_surv[pclass_surv['Fare'] > 20]
rich_pass
```

```
[57]:
```

	Fare	Survived	Pclass
1	71.2833	1	1
3	53.1000	1	1
6	51.8625	0	1
7	21.0750	0	3
9	30.0708	1	2
..
880	26.0000	1	2
885	29.1250	0	3
887	30.0000	1	1
888	23.4500	0	3
889	30.0000	1	1

[376 rows x 3 columns]

Now fetch the same for passengers that paid 20 pounds or less for their voyage.

```
[59]: # passenger class and survival of low fares
poor_pass = pclass_surv[pclass_surv['Fare'] <= 20]
poor_pass
```

```
[59]:
```

	Fare	Survived	Pclass
0	7.2500	0	3
2	7.9250	1	3
4	8.0500	0	3
5	8.4583	0	3
8	11.1333	1	3
..
882	10.5167	0	3
883	10.5000	0	2
884	7.0500	0	3
886	13.0000	0	2
890	7.7500	0	3

[515 rows x 3 columns]

Finally, get the class and survival status for passengers that paid either less than 10 lbs. **or** more than 50 lbs. for their fare.

```
[63]: # ppl paying a little *or* a lot
extremes_pass = pclass_surv[(pclass_surv['Fare'] < 50) | (pclass_surv['Fare'] > 50)]
extremes_pass
```

```
[63]:
```

	Fare	Survived	Pclass
0	7.2500	0	3
1	71.2833	1	1
2	7.9250	1	3
3	53.1000	1	1
4	8.0500	0	3
...
886	13.0000	0	2
887	30.0000	1	1
888	23.4500	0	3
889	30.0000	1	1
890	7.7500	0	3

[891 rows x 3 columns]

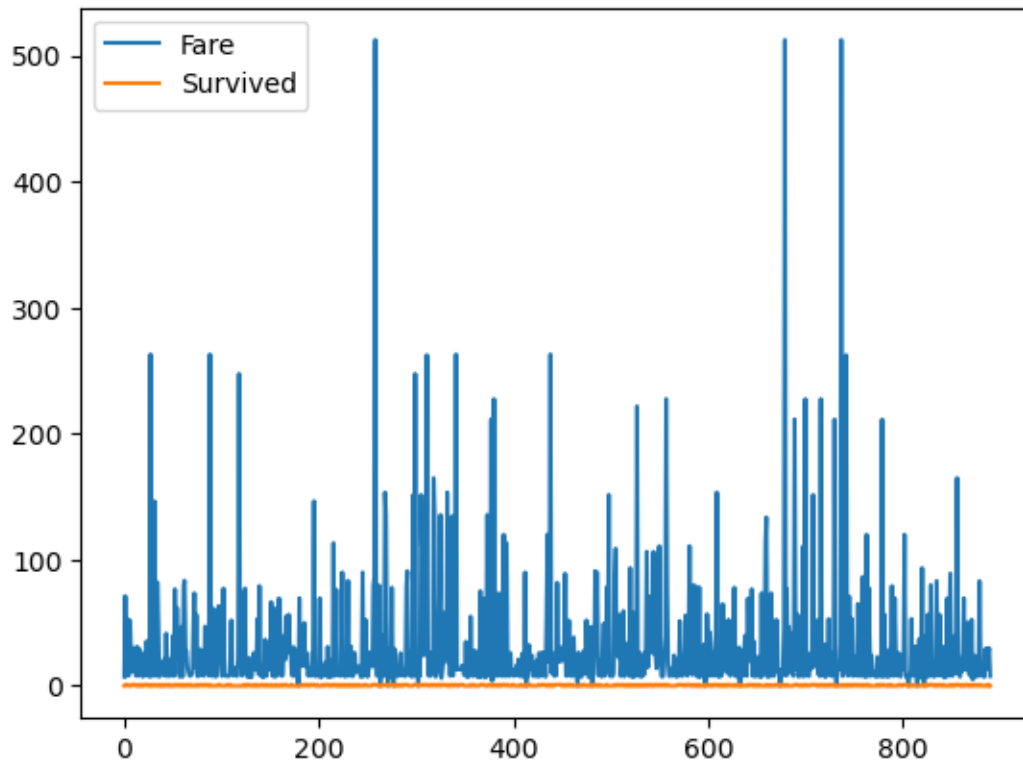
If you did the above in two steps, see if you can do it in one go instead! There are hints just above.

1.4 Basic Plotting

`DataFrame` objects know how to plot themselves! Or, more precisely, `DataFrame` objects have methods for plotting. Let's try!

```
[64]: import matplotlib as plt
fare_surv.plot()
```

```
[64]: <AxesSubplot:>
```



As a graph, this one isn't very informative, but it does show us what the default `DataFrame.plot()` method does: it plots (numerical) data by row index. This could be quite useful if a data frame were sorted on a particular variable...

Other type of plots are reached through plot, like `fare_surv.plot.scatter()` or similar. We can see what methods are available by hitting the `<TAB>` key after `DataFrame.plot.`

Do this below;

```
[77]: fare_surv.plot.
```

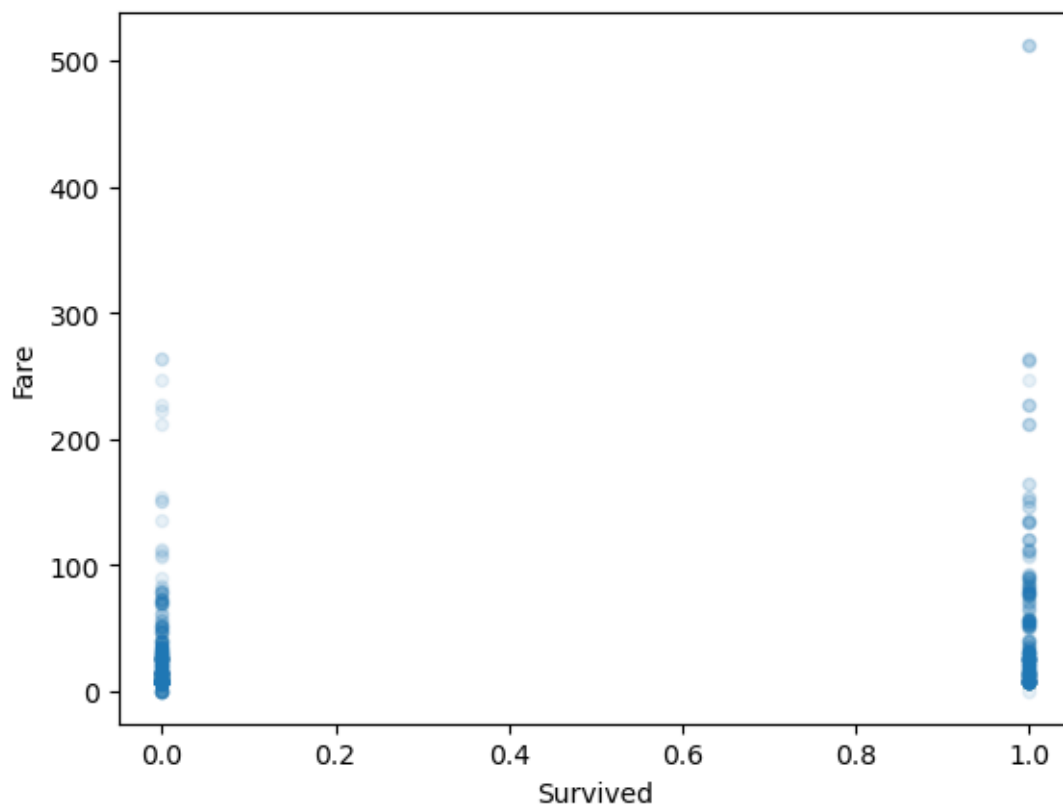
```
[77]: <pandas.plotting._core.PlotAccessor object at 0x7fcb083aaa30>
```

So there *is* a `scatter()` available, along with many of our other `matplotlib` friends.

Let's try a scatter plot Fare vs. Survival.

```
[78]: fare_surv.plot.scatter(x="Survived", y="Fare", alpha = 0.1)
```

```
[78]: <AxesSubplot:xlabel='Survived', ylabel='Fare'>
```

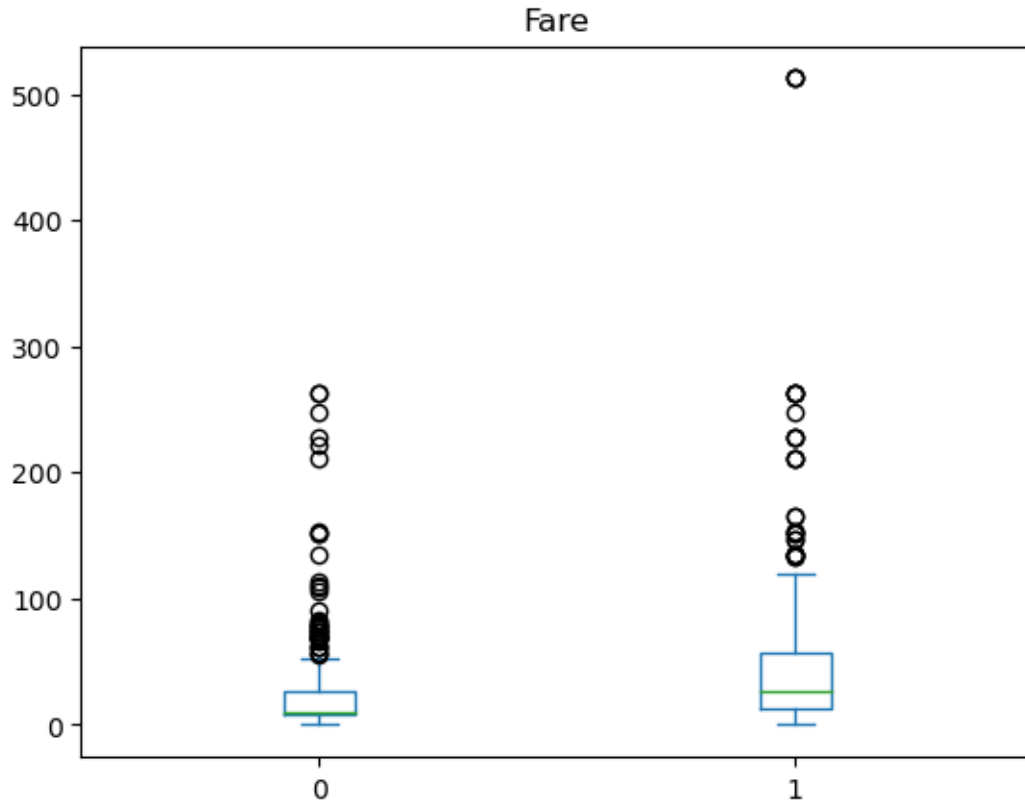


Looks like those 500 lb. fares were worth it.

Use the cell below to make a box plot of the *column Fare by* the variable **Survived**.

```
[82]: # boxplot of Fare paid by Survival status
fare_surv.plot.box(column = "Fare", by = "Survived")
```

```
[82]: Fare      AxesSubplot(0.125,0.11;0.775x0.77)
dtype: object
```



1.5 Calculating New Columns

We often want to compute new columns based on existing ones. Pandas makes this really easy!

Let's use numpy to make a toy data set of annual wages and income-from-interest for 10 people.

```
[83]: import numpy as np
```

In following code, you should be able to understand the numpy bit up top. The pandas bit further down should sort of make sense, but don't worry if you don't fully understand it. You can come back and look at it again after you've finished this tutorial.

```
[84]: # make some incomes in thousands of US dollars
rng = np.random.default_rng(seed=42)
raw_dat = rng.integers(0,100,size=(10, 2))
raw_dat[:,0] = raw_dat[:,0] + 100
raw_dat[4,1] = raw_dat[4,1] + 200

# make initial column names
col_names = ['wage', 'interest']
```

```
# make the initial pandas data frame
incomes = pd.DataFrame(raw_dat, columns = col_names)

# add a gender column
gender = ['f', 'm', 'n', 'f', 'f', 'n', 'm', 'm', 'f', 'f']
incomes['gender'] = gender

# look at our new data frame
incomes
```

```
[84]:
```

	wage	interest	gender
0	108	77	f
1	165	43	m
2	143	85	n
3	108	69	f
4	120	209	f
5	152	97	n
6	173	76	m
7	171	78	m
8	151	12	f
9	183	45	f

One obvious thing to look at from a behavioral science perspective would be total income. After all, money is money...

So we'll make a new column for total income, and set it to the sum of the wage and interest columns. To do this, we address our desired column as though it already exists, and make it equal to what we want (the sum of wage and interest income, in this case).

```
[85]: incomes['total'] = incomes['wage'] + incomes['interest']
incomes
```

```
[85]:
```

	wage	interest	gender	total
0	108	77	f	185
1	165	43	m	208
2	143	85	n	228
3	108	69	f	177
4	120	209	f	329
5	152	97	n	249
6	173	76	m	249
7	171	78	m	249
8	151	12	f	163
9	183	45	f	228

All of the *arithmetic* and *logical* operators can be used to create new columns based on existing ones.

We can also use scalar multipliers or addends, etc. (like we did when we created the raw data with

numpy just above). The scaler will be “broadcast” to each element of the column.

For example, if we wanted to know the total income in Euros, we could do this:

```
[86]: dol2eu = 0.94 # 0.94 euros per US dollar (early 2023)
      incomes['total_eu'] = dol2eu * incomes['total']
      incomes
```

```
[86]:
```

	wage	interest	gender	total	total_eu
0	108	77	f	185	173.90
1	165	43	m	208	195.52
2	143	85	n	228	214.32
3	108	69	f	177	166.38
4	120	209	f	329	309.26
5	152	97	n	249	234.06
6	173	76	m	249	234.06
7	171	78	m	249	234.06
8	151	12	f	163	153.22
9	183	45	f	228	214.32

In the cell below, add a Boolean (True/False) column that shows if each person’s wages exceeds their income from interest.

```
[96]: # adding a wages vs incomes comparison column
      incomes['wage_vs_income'] = incomes['total'] > incomes['wage']
      incomes
```

```
[96]:
```

	wage	interest	gender	total	total_eu	wage_vs_income
0	108	77	f	185	173.90	True
1	165	43	m	208	195.52	True
2	143	85	n	228	214.32	True
3	108	69	f	177	166.38	True
4	120	209	f	329	309.26	True
5	152	97	n	249	234.06	True
6	173	76	m	249	234.06	True
7	171	78	m	249	234.06	True
8	151	12	f	163	153.22	True
9	183	45	f	228	214.32	True

1.6 Summary Statistics

Getting summary statistics is also something that pandas makes really easy.

1.6.1 Simple descriptive statistics

We can get a quick look an entire `DataFrame` with its `describe()` method (similar to `summary()` in R).

```
[97]: incomes.describe()
```

```
[97]:
```

	wage	interest	total	total_eu
count	10.000000	10.000000	10.000000	10.000000
mean	147.400000	79.100000	226.500000	212.910000
std	27.281251	51.977452	47.815037	44.946135
min	108.000000	12.000000	163.000000	153.220000
25%	125.750000	51.000000	190.750000	179.305000
50%	151.500000	76.500000	228.000000	214.320000
75%	169.500000	83.250000	249.000000	234.060000
max	183.000000	209.000000	329.000000	309.260000

Notice that `describe()` handled the presence of a string column gracefully by ignoring it rather than producing an error.

If we hit the <TAB> key after `incomes.`, we'll see that `DataFrame` objects have a LOT of methods!

```
[98]: incomes.
```

```
File "/var/folders/zc/6v283x0929j5f38j6cvlwbwr0000gn/T/ipykernel_83524/
↪2418687937.py", line 1
    incomes.
    ~
SyntaxError: invalid syntax
```

If we browse around a little, we see that all the common summary statistics like mean, median, standard deviation, etc. are there, and they all have reasonable names. Let's compute the mean

```
[99]: incomes.mean()
```

```
/var/folders/zc/6v283x0929j5f38j6cvlwbwr0000gn/T/ipykernel_83524/1563429750.py:1
: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with
'numeric_only=None') is deprecated; in a future version this will raise
TypeError. Select only valid columns before calling the reduction.
    incomes.mean()
```

```
[99]: wage                147.40
      interest            79.10
      total              226.50
      total_eu           212.91
      wage_vs_income      1.00
      dtype: float64
```

That worked, but it complained (at least my version of pandas did). It wants us to pick only valid (numeric) columns over which to compute the mean. Okay.

```
[100]: incomes[['wage', 'interest']].mean()
```

```
[100]: wage          147.4
      interest       79.1
      dtype: float64
```

Compute the standard deviation of total income (in Euros, if you prefer)

```
[104]: # deviation of total income
      incomes[['total_eu']].std()
```

```
[104]: total_eu      44.946135
      dtype: float64
```

Pro tip: if you *do* want to compute a statistic on *all* the numeric columns on large data frame, you can save typing with `DataFrame.mean(numeric_only = True)`. Try it!

```
[103]: incomes.mean(numeric_only = True)
```

```
[103]: wage          147.40
      interest       79.10
      total         226.50
      total_eu       212.91
      wage_vs_income 1.00
      dtype: float64
```

1.6.2 Computing statistics by group

We can also easily compute statistics separately based on a grouping variable, like ‘gender’ for the incomes data.

Here’s our grouping variable:

```
[105]: incomes['gender']
```

```
[105]: 0    f
      1    m
      2    n
      3    f
      4    f
      5    n
      6    m
      7    m
```

```
8    f
9    f
Name: gender, dtype: object
```

And now we'll use it in our data frame's `groupby()` method. Like this.

```
[106]: incomes[['total', 'gender']].groupby('gender').mean()
```

```
[106]:          total
gender
f      216.400000
m      235.333333
n      238.500000
```

If you are coming from the R/tidyverse world (e.g. if you took PSY420 recently), you'll recognize this command as similar to using the pipe (`%>%`).

What's happening is that

- `incomes[['total', 'gender']]` creates a data frame
- `groupby('gender')` creates another data frame grouped by gender
- `mean()` computes the mean on the grouped data frame

So we could (almost) turn this directly into R code that uses the pipe:

```
incomes[['total', 'gender']] %>%
groupby('gender') %>%
mean()
```

How many people were in each group? Just use the `value_counts()` method!

```
[107]: incomes['gender'].value_counts()
```

```
[107]: f      5
      m      3
      n      2
Name: gender, dtype: int64
```

In the cell below, compute the survival rate for passengers on the RMS Titanic grouped by passenger class.

(*hint* - having the Survived variable coded as 0 or 1 works to your advantage)

```
[110]: titanic[['Pclass', 'Survived']].groupby('Pclass').mean()
```

```
[110]:          Survived
Pclass
1      0.629630
2      0.472826
```

1.6.3 Multiple statistics using aggregation

We can compute many things at once using the `agg()` (aggregate) method. To use this method, we pass it a dictionary in which the keys are column names and the values are lists of valid statistics (i.e. methods that `DataFrames` know about). Like this.

```
[111]: my_stats_dict = {  
        "wage": ["mean", "std"],  
        "interest": ["mean", "std"],  
        "total": ["mean", "std"]  
    }  
  
    incomes.agg(my_stats_dict)
```

```
[111]:
```

	wage	interest	total
mean	147.400000	79.100000	226.500000
std	27.281251	51.977452	47.815037

You can do the above in one go (rather than defining a separate `my_stats_dict` object), but it looks a bit messy in our opinion.